

```

!pip install seqeval --quiet
----- 43.6/43.6 kB 1.3 MB/s eta
0:00:00
  etadata (setup.py) ...

import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torch.nn.utils.rnn import pad_sequence
from datasets import load_dataset
from tqdm import tqdm
from seqeval.metrics import classification_report, f1_score

BATCH_SIZE = 16
EMBEDDING_DIM = 100
HIDDEN_DIM = 256
OUTPUT_SIZE = 9
NUM_LAYERS = 1
EPOCHS = 10
LEARNING_RATE = 0.001
PAD_TAG = -1
PATIENCE = 3

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")

Using device: cpu

```

## Task 1

```

dataset = load_dataset("lhoestq/conll2003")
train_data = dataset["train"]
val_data = dataset["validation"]
test_data = dataset["test"]

{"model_id": "62b4f9e67b6d4cfaa7bd62303580b283", "version_major": 2, "version_minor": 0}

{"model_id": "b98aab4a0cc48bf8e8bc0e381c9fc8f", "version_major": 2, "version_minor": 0}

{"model_id": "635c0bb31fb4833bc9d0f77c1da0c12", "version_major": 2, "version_minor": 0}

{"model_id": "d74d2adc9d944f9f888358917bc44d78", "version_major": 2, "version_minor": 0}

```

```

{"model_id": "9623e26074824c4c864e888c0ab38973", "version_major": 2, "version_minor": 0}

{"model_id": "e4d8c0c751dc4112817e70cad06818cd", "version_major": 2, "version_minor": 0}

{"model_id": "e5568ac48e4e496784d1812ae53fa8be", "version_major": 2, "version_minor": 0}

word_to_ix = {"<PAD>": 0, "<UNK>": 1}
idx = 2
for sent in train_data["tokens"]:
    for tok in sent:
        if tok not in word_to_ix:
            word_to_ix[tok] = idx
        idx += 1

print(f"Vocab size: {len(word_to_ix)})")

id2tag = [
    "O", "B-PER", "I-PER", "B-ORG", "I-ORG",
    "B-LOC", "I-LOC", "B-MISC", "I-MISC"
]
Vocab size: 23625

```

## Task 2

```

class NERDataset(Dataset):
    def __init__(self, sentences, tags_ids, word_to_ix):
        self.sentences = sentences
        self.tags_ids = tags_ids
        self.word_to_ix = word_to_ix
        self.unk_id = word_to_ix.get("<UNK>", 1)

    def __len__(self):
        return len(self.sentences)

    def __getitem__(self, idx):
        sent = self.sentences[idx]
        tag_ids = self.tags_ids[idx]

        sent_ids = [self.word_to_ix.get(tok, self.unk_id) for tok in sent]

        return torch.tensor(sent_ids, dtype=torch.long),
               torch.tensor(tag_ids, dtype=torch.long)

    def collate_fn(batch):

```

```

        sentences, tags = zip(*batch)
        padded_sentences = pad_sequence(sentences, batch_first=True,
padding_value=0)
        padded_tags = pad_sequence(tags, batch_first=True,
padding_value=PAD_TAG)
        return padded_sentences, padded_tags

train_loader = DataLoader(
    NERDataset(train_data["tokens"], train_data["ner_tags"],
word_to_ix),
    batch_size=BATCH_SIZE, shuffle=True, collate_fn=collate_fn
)

val_loader = DataLoader(
    NERDataset(val_data["tokens"], val_data["ner_tags"], word_to_ix),
    batch_size=BATCH_SIZE, shuffle=False, collate_fn=collate_fn
)

test_loader = DataLoader(
    NERDataset(test_data["tokens"], test_data["ner_tags"],
word_to_ix),
    batch_size=BATCH_SIZE, shuffle=False, collate_fn=collate_fn
)

```

## Task 3

```

class NERLSTMModel(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim,
output_size, num_layers=1):
        super(NERLSTMModel, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.lstm = nn.LSTM(
            input_size=embedding_dim,
            hidden_size=hidden_dim,
            num_layers=num_layers,
            batch_first=True,
            bidirectional=True)
        self.fc = nn.Linear(hidden_dim * 2, output_size)

    def forward(self, x):
        embedded = self.embedding(x)
        lstm_out, _ = self.lstm(embedded)
        logits = self.fc(lstm_out)
        return logits

```

## Task 4

```
model = NERLSTMModel(len(word_to_ix), EMBEDDING_DIM, HIDDEN_DIM,
OUTPUT_SIZE, NUM_LAYERS).to(device)
criterion = nn.CrossEntropyLoss(ignore_index=PAD_TAG)
optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)

trigger_times = 0
best_val_loss = float('inf')

for epoch in range(EPOCHS):
    model.train()
    total_loss = 0.0

    progress_bar = tqdm(train_loader, desc=f"Epoch
{epoch+1}/{EPOCHS}")

    for sentences, tags in progress_bar:
        sentences, tags = sentences.to(device), tags.to(device)

        optimizer.zero_grad()
        predictions = model(sentences)

        predictions = predictions.view(-1, OUTPUT_SIZE)
        tags = tags.view(-1)

        loss = criterion(predictions, tags)
        loss.backward()
        optimizer.step()

        total_loss += loss.item()
        progress_bar.set_postfix(loss=loss.item())

    avg_train_loss = total_loss / len(train_loader)

    model.eval()
    total_val_loss = 0.0
    val_correct = 0
    val_total = 0

    with torch.no_grad():
        for sentences, tags in val_loader:
            sentences, tags = sentences.to(device), tags.to(device)

            outputs = model(sentences) # [Batch, Seq, Output]

            outputs_flat = outputs.view(-1, OUTPUT_SIZE)
            tags_flat = tags.view(-1)
            loss = criterion(outputs_flat, tags_flat)
            total_val_loss += loss.item()
```

```
preds = torch.argmax(outputs, dim=2)
mask = (tags != PAD_TAG) # Bo' qua padding
val_correct += (preds[mask] == tags[mask]).sum().item()
val_total += tags[mask].numel()

avg_val_loss = total_val_loss / len(val_loader)
val_acc = val_correct / val_total if val_total > 0 else 0

print(f"Epoch {epoch+1}: Train Loss: {avg_train_loss:.4f} | Val Loss: {avg_val_loss:.4f} | Val Acc: {val_acc:.4f}")

if avg_val_loss < best_val_loss:
    best_val_loss = avg_val_loss
    trigger_times = 0
else:
    trigger_times += 1

if trigger_times >= PATIENCE:
    print("Early Stopping! Stop training at epoch {epoch+1}")
    break

Epoch 1/10: 100%|██████████| 878/878 [00:59<00:00, 14.72it/s, loss=0.126]
Epoch 1: Train Loss: 0.3705 | Val Loss: 0.2642 | Val Acc: 0.9262
Epoch 2/10: 100%|██████████| 878/878 [01:01<00:00, 14.36it/s, loss=0.114]
Epoch 2: Train Loss: 0.1588 | Val Loss: 0.2051 | Val Acc: 0.9431
Epoch 3/10: 100%|██████████| 878/878 [01:01<00:00, 14.26it/s, loss=0.0747]
Epoch 3: Train Loss: 0.0716 | Val Loss: 0.1841 | Val Acc: 0.9483
Epoch 4/10: 100%|██████████| 878/878 [01:01<00:00, 14.29it/s, loss=0.0154]
Epoch 4: Train Loss: 0.0270 | Val Loss: 0.1942 | Val Acc: 0.9519
Epoch 5/10: 100%|██████████| 878/878 [01:02<00:00, 14.13it/s, loss=0.00188]
Epoch 5: Train Loss: 0.0089 | Val Loss: 0.2128 | Val Acc: 0.9525
Epoch 6/10: 100%|██████████| 878/878 [01:02<00:00, 14.13it/s, loss=0.00164]
Epoch 6: Train Loss: 0.0035 | Val Loss: 0.2415 | Val Acc: 0.9526
Early Stopping! Stop training at epoch 6
```

## Task 5

```
def predict_sentence(sentence):
    model.eval()
    tokens = sentence.split()
    unk_id = word_to_ix.get("<UNK>", 1)
    input_ids = [word_to_ix.get(word, unk_id) for word in tokens]
    input_tensor = torch.tensor([input_ids],
                               dtype=torch.long).to(device)

    with torch.no_grad():
        outputs = model(input_tensor)
        prediction_indices = torch.argmax(outputs, dim=2)

    predicted_tags = [id2tag[idx] for idx in
                      prediction_indices[0].tolist()]

    print(f"\nResult for: '{sentence}'")
    print(f"{'Word':<15} {'Pred Tag':<15}")
    print("-" * 30)
    for word, tag in zip(tokens, predicted_tags):
        print(f"{word:<15} {tag:<15}")

test_sentence = "On Monday, October 14, 2024, Google CEO Sundar Pichai traveled from Mountain View, California to New York City to meet with Microsoft executives including Satya Nadella and Amy Hood at the headquarters of the United Nations, where they announced a joint initiative with NASA and SpaceX to develop advanced AI-powered satellites capable of monitoring climate change, while also signing agreements with the European Space Agency, Tesla, and Siemens to collaborate on renewable energy projects that aim to reduce global carbon emissions by 30 percent within the next decade, all before attending a high-profile conference at Harvard University hosted by the World Economic Forum, which included keynote speeches from former U.S. President Barack Obama, Nobel laureate Malala Yousafzai, and philanthropist Bill Gates."
predict_sentence(test_sentence)
```

```
Result for: 'On Monday, October 14, 2024, Google CEO Sundar Pichai traveled from Mountain View, California to New York City to meet with Microsoft executives including Satya Nadella and Amy Hood at the headquarters of the United Nations, where they announced a joint initiative with NASA and SpaceX to develop advanced AI-powered satellites capable of monitoring climate change, while also signing agreements with the European Space Agency, Tesla, and Siemens to collaborate on renewable energy projects that aim to reduce global carbon emissions by 30 percent within the next decade, all before attending a high-profile conference at Harvard University hosted by the World Economic Forum, which included keynote speeches from former
```

U.S. President Barack Obama, Nobel laureate Malala Yousafzai, and philanthropist Bill Gates.'

Word	Pred	Tag
------	------	-----

Word	Pred	Tag
On	0	
Monday,	0	
October	0	
14,	B-MISC	
2024,	0	
Google	B-MISC	
CEO	0	
Sundar	0	
Pichai	B-MISC	
traveled	0	
from	0	
Mountain	B-LOC	
View,	0	
California	B-LOC	
to	0	
New	B-LOC	
York	I-LOC	
City	I-LOC	
to	0	
meet	0	
with	0	
Microsoft	B-ORG	
executives	0	
including	0	
Satya	B-ORG	
Nadella	I-PER	
and	0	
Amy	B-PER	
Hood	0	
at	0	
the	0	
headquarters	0	
of	0	
the	0	
United	B-LOC	
Nations,	0	
where	0	
they	0	
announced	0	
a	0	
joint	0	
initiative	0	
with	0	
NASA	0	
and	0	

SpaceX 0  
to 0  
develop 0  
advanced 0  
AI-powered 0  
satellites 0  
capable 0  
of 0  
monitoring 0  
climate 0  
change, 0  
while 0  
also 0  
signing 0  
agreements 0  
with 0  
the 0  
European B-MISC  
Space I-MISC  
Agency, I-MISC  
Tesla, I-MISC  
and 0  
Siemens B-MISC  
to 0  
collaborate 0  
on 0  
renewable 0  
energy 0  
projects 0  
that 0  
aim 0  
to 0  
reduce 0  
global 0  
carbon 0  
emissions 0  
by 0  
30 0  
percent 0  
within 0  
the 0  
next 0  
decade, 0  
all 0  
before 0  
attending 0  
a 0  
high-profile 0  
conference 0

```

at          0
Harvard    B-ORG
University I-ORG
hosted     0
by         0
the        0
World      B-MISC
Economic   I-MISC
Forum,     I-MISC
which      0
included   0
keynote   0
speeches   0
from       0
former     0
U.S.       B-LOC
President  0
Barack     B-PER
Obama,     I-PER
Nobel      I-PER
laureate   I-PER
Malala     0
Yousafzai, I-PER
and        0
philanthropist B-MISC
Bill       I-PER
Gates.    I-PER

def align_predictions(predictions, label_ids, id2tag, pad_tag_id=-1):
    preds_list = []
    labels_list = []
    for i in range(len(label_ids)):
        sent_preds = []
        sent_labels = []
        for j in range(len(label_ids[i])):
            if label_ids[i][j] != pad_tag_id:
                sent_preds.append(id2tag[predictions[i][j]])
                sent_labels.append(id2tag[label_ids[i][j]])
        preds_list.append(sent_preds)
        labels_list.append(sent_labels)
    return preds_list, labels_list

def evaluate_ner_metrics(model, data_loader, id2tag, device,
                        pad_tag_id=-1):
    model.eval()
    all_preds = []
    all_labels = []
    with torch.no_grad():
        for sentences, tags in data_loader:
            sentences = sentences.to(device)

```

```

tags = tags.to(device)
outputs = model(sentences)
predictions = torch.argmax(outputs, dim=2)
predictions_np = predictions.cpu().numpy()
tags_np = tags.cpu().numpy()
batch_preds, batch_labels =
align_predictions(predictions_np, tags_np, id2tag, pad_tag_id)
all_preds.extend(batch_preds)
all_labels.extend(batch_labels)
print("\n" + "*"*55)
print(classification_report(all_labels, all_preds))
print(f"Macro F1-Score: {f1_score(all_labels, all_preds):.4f}")
print("*"*55 + "\n")

def evaluate(model, data_loader):
    model.eval()
    correct_preds = 0
    total_preds = 0
    with torch.no_grad():
        for sentences, tags in data_loader:
            sentences, tags = sentences.to(device), tags.to(device)
            outputs = model(sentences)
            predictions = torch.argmax(outputs, dim=2)

            mask = (tags != PAD_TAG)
            valid_predictions = predictions[mask]
            valid_tags = tags[mask]

            correct_preds += (valid_predictions ==
valid_tags).sum().item()
            total_preds += valid_tags.numel()

    return correct_preds / total_preds if total_preds > 0 else 0.0

evaluate(model, test_loader)
0.9307418972757618

evaluate_ner_metrics(model, test_loader, id2tag, device)

```

	precision	recall	f1-score	support
LOC	0.86	0.68	0.76	1668
MISC	0.45	0.64	0.53	702
ORG	0.74	0.57	0.64	1661
PER	0.65	0.67	0.66	1617
micro avg	0.69	0.64	0.66	5648
macro avg	0.68	0.64	0.65	5648

weighted avg	0.71	0.64	0.67	5648
--------------	------	------	------	------

Macro F1-Score: 0.6622

---